



Medical Image Synthesis and Augmentation Using GANs
(Focus on MedGAN and Beyond)

□ Student Name- Simranjeet Kaur

□ Course Name – ADVANCED MACHINE LEARNING

□ Instructor Name - Chaojiang (CJ) Wu

Spring 2025

Table of Contents

- 1. Introduction**
- 2. Literature Review**
- 3. Industry Applications**
- 4. Future Directions and Challenges**
- 5. Conclusion**
- 6. References**

1. Introduction

The incorporation of artificial intelligence in health care has made significant changes in the traditional medical practice in terms of improving the correctness of patient care, working faster, and increasing accessibility. Among the different AI methodologies, deep learning stands out as perhaps the most pursued technology due to its capacity for handling large volumes of data, unearthing complex patterns, and making highly accurate predictions. Today, deep learning has cemented its place in the enhancement of health care through medical imaging, disease identification, and the development of individualized treatment plans [1].

With the dawning of this new era of health care, the value of comprehensive medical data also is rising. In this context, healthcare accuracy refers to all patient information, ranging from molecular characteristics, EHRs, and environmental factors to lifestyle. Thus, the right drug should go to the right patient at the right time. With more access to healthcare data, new opportunities and challenges arise in research. Unearthing the relationships in massive medical databases important stumbling block in designing a robust, data-driven healthcare system-is one of the challenges arising alongside the new opportunities. While deep learning and machine learning techniques for prediction and discovery hold great promise, they have thus far failed to gain much purchase in the field of health, owing to several strong competing factors. Particularly, the complexity, heterogeneity, temporal dependence, sparsity, and stochastic nature of biological data comparatively hinder their efficient applications [2].




Figure 1: A representation of various applications of AI in healthcare [1]

Top-Notch Understanding Over the Last Ten Years: Machine Learning Has Turned from a Technology into an Everyday Fact-of-Life Application, Mainly because of Deep Neural Networks [3]. They Perform Exceptionally Well on Various Clinical Imaging Tasks. Nowadays, the Birth of Generative Adversarial Networks (GANs) has had a HUGE Impact on Generative Modeling and Data Synthesis, Leveling Off the Quality Achieved with Such Techniques. The Speed at which Research on GANs Has Continued Intensifying Has Been Endlessly Pushed by the Bar Raising Levels of Image Quality. A Landmark in the Field Came with BigGAN [4], Which Gave an Amazing Lift to Image Generation. StyleGAN [5] then Followed Suit with another Major Jump, Producing Extremely Realistic High-Resolution Images of Human Faces. Inspired by the Historical Gold Standards for Well-Performing GANs in Natural Images, This Study Aims to Measure Efficacy in the Application of GANs to Medical Datasets, Often Small and Tight with Anatomical Constraints. Even though several of the reviews have explored the use of GANs in medical imaging [6][7], this work stands out in its empirical demonstration of the advantages of data generated through GANs, combined with a thorough hyperparameter review of the different methods.

In recent developments, GANs have gained considerable attention from the medical research community because of their potential for creating realistic medical images. For instance, [8] proposed a GAN model designed to generate synthetic T1-weighted brain MRIs-based scans closely reflecting original ones [9]. They also obtained high-resolution images of skin lesions that the experts could hardly distinguish between the real ones and synthetic ones. used GANs to produce brain MRIs and found that these belonged to high performance in qualitative and quantitative evaluations. Furthermore, in [10], it was revealed that GAN-generated lung cancer nodule images are virtually indistinguishable from actual ones, even by experienced radiologists.

2. Literature Review

GANs have also been investigated to generate additional training data. In [11], a GAN was trained to generate synthetic brain tumor MRIs, and the effect of using these data to train segmentation networks was studied. The results showed that network segmentation models trained only with synthetic data did not perform as well as those trained with real data. In the same manner, [12] presents a framework that integrates a variational autoencoder with a GAN used for data augmentation on image segmentation tasks. Even in this case, the method had some potential; however, the results remained mixed as to how effective it was. As stated in the review [13], GANs are now employed in medical imaging not only for image generation but also for domain adaptation, classification, image reconstruction, and so on. It is concerning, however, that despite the ability of GANs to generate plausible images, there remain doubts about their genuine clinical worth in comparison with real medical data when applied to day-to-day diagnostic tasks.

They are trained on data up to October 2021, but can do some projections for data beyond this time. Medical imaging has specific important aspects in healthcare; however, the small size of medical datasets usually limits the development of algorithms. This paper reviewed numerous GAN architecture types that have been used for generating synthetic medical images, augmenting training data. Although they produce excellent-looking images, the generated images do not perform as well as real data in specific applications; this indicates much caution in their use and the need for tailoring particular solutions of GANs [14]. Medical Image Synthesis (MIS) is very important in reducing diagnostic costs, but it is fraught with problems in maintaining biological relevance on account of the complexity of medical images.

To tackle this, we proposed the Hybrid Augmented GAN (HAGAN), which augments the AttnMix Generator with a Hierarchical Discriminator and Reverse Skip Connection for improving structural realism and fine details. The experiments conducted on COVID-CT, ACDC, and BraTS2018 datasets show that HAGAN outperforms contemporary methods and achieves state-of-the-art results at all resolutions [15].




Figure 2: Flowchart of GAN [16]

Simply put, GANs (adversarial networks) see liabilities between a generator and a discriminator playing a game of minimax: the generator parameters θ generate as realistic samples as possible from random noise, whereas the discriminator parameters ϕ distinguish real data from fake data. Both networks are trained at

the same time using an adversarial loss [16] that induces the generator to improve its output to deceive the discriminator more and more.

Table 1: the key points from the literature review on GANs in medical imaging

Study/Framework	Purpose/Focus	Key Findings/Conclusions
[10] GAN for Synthetic Brain Tumor MRI Generation	Investigated GAN-generated synthetic brain tumor MRIs for training segmentation networks.	Models trained with synthetic data alone performed worse than those trained with real data, highlighting limitations in GAN use.
[11] GAN + Variational Autoencoder for Data Augmentation	A framework combining GAN with a variational autoencoder for data augmentation in image segmentation tasks.	Results were mixed regarding effectiveness; while promising, further refinement is needed.
[12] GANs in Medical Imaging	Reviewed GAN applications in medical imaging for tasks such as image generation, domain adaptation, and image reconstruction.	Despite generating plausible images, GANs may not be as reliable as real medical data for clinical use.
[13] GAN Architecture for Synthetic Medical Image Generation	Reviewed various GAN architectures for generating synthetic medical images to augment training data.	While images generated by GANs look realistic, their performance in specific applications lags behind real data, necessitating caution.
[14] Hybrid Augmented GAN (HAGAN) for Medical Image Synthesis	Proposed HAGAN model for improving structural realism and fine details in medical image synthesis.	HAGAN outperformed other methods, achieving state-of-the-art results on multiple medical image datasets (COVID-CT, ACDC, BraTS2018).

6

3. Industry Applications

Diagnosis of rare diseases takes an average of five years, and nearly 400 million people are affected by rare conditions worldwide. However, most of the patients go undiagnosed, or in many cases, they are misdiagnosed. The present study aims to employ ChatGPT-based Retrieval Augmented Generation (RAG) in rare disease diagnostic assistance. The RareDxGPT model, which takes information from the RareDis Corpus, could then achieve success in rare disease diagnosis at 40% accuracy using a simple prompt. More detailed prompts improved accuracy, "Prompt + Explanation," reached 43% accuracy for RareDxGPT. These early results

indicate the potential for the application of AI ameliorated with domain knowledge in assisting rare disease diagnosis among other areas [17].

This article presents the MRI-to-CT transformer-based improved denoising diffusion probabilistic model (MC-IDDPM) for synthesizing high-quality synthetic computed tomographic images (sCTs) from MRIs, thus making it radiation therapy planning more efficient by eliminating the use of CT simulation and reducing image registration errors. The model employs a shifted-window transformer network to add Gaussian noise to CT scans and

7

Then denoise them using a Swin-Vnet conditioned on MRI parameters. Evaluated on brain and prostate datasets, MC-IDDPM produced some very impressive results, with miles ahead of the pack on state-of-the-art metrics such as MAE, PSNR, SSIM, and NCC, and significant statistical improvements over alternative approaches. In addition, dosimetric analyses indicate relatively minor differences in target dose coverage, thus emphasizing the capability of the model in optimizing the entire resource-reducing exposure of patients to radiation [18]. Generative adversarial networks (GANs) have emerged as an attractive solution, enhancing image quality and minimizing radiation exposure in low-dose computed tomography (LDCT) imaging. This review focuses on the development of GAN-based LDCT denoising techniques using several architectures including cGAN, CycleGAN, and super-resolution GAN (SRGAN). Their performances were assessed using qualitative as well as quantitative metrics, i.e., PSNR, SSIM, and LPIPS, and importantly, substantial improvements were found in both benchmark and clinical datasets. Nevertheless, while promising, this approach faces hurdles such as the interpretability of images, synthetic artifacts, and the lack of clinically relevant metrics for evaluation. Overall, they have the potential to impact precision medicine in radiology; however, these models are required to evolve further to allow for broader adoption in the clinical space [19].

The main objective of this study is to bridge the gap that now exists between personalized medicine and healthcare by merging the involvement of forecasting diagnosis tracks and individualized treatment through artificial intelligence (AI). With further developments in big data analytics, machine learning, and deep neural networks, this research proposes to explore the impact AI has on reshaping healthcare through precise prediction for disease risk assertions, diagnosis, and treatment responses for each patient. For instance, AI includes using deep neural networks for skin cancer classification at the level of expertise of a dermatologist and, even in cardiology, for risk predictions and diagnosis. This study reviews the existing literature and the recently challenged issues of technology, providing insight into the opportunities AI-based

Personalized medicine will help humans to improve their health outcomes, reduce waste of resources, and ultimately minimize health equity issues. Second, this paper examines the implications of personalized medicine approaches on the US healthcare system regarding chronic conditions like diabetes. These days, personalized treatment plans based on the patient profile are favored over one medication regimen for all patients. It focuses, similarly, on productivity improvement, cost reduction, and patient-centered treatment. It finally sketches the projection of personalized medicine forward in the US on two main topics: current challenges and opportunities for its further development and adoption as the integration of genomic data into clinical decision-making to fit treatment for cancer patients [20].

Table 2: A summary of the industry applications you provided, organized into a table format:

Industry Application	Description	Key Outcomes/Implications
Rare Disease Diagnosis with ChatGPT-based RAG (RareDxGPT)	ChatGPT-based Retrieval Augmented Generation (RAG) model to assist in rare disease diagnosis using the RareDis Corpus.	40% accuracy with simple prompts, improved to 43% with more detailed "Prompt + Explanation" approach, highlighting the potential of AI and domain knowledge in rare disease diagnosis.
MRI-to-CT Transformer-based MC-IDDPM for Synthetic CT Image Generation	MC-IDDPM model synthesizes high-quality synthetic CT images from MRIs using a shiftedwindow transformer network.	Achieved significant improvements in MAE, PSNR, SSIM, and NCC, enabling more efficient radiation therapy planning by reducing radiation exposure and image registration errors.
Generative Adversarial Networks (GANs) for LDCT Denoising	GAN-based denoising techniques for improving low-dose computed tomography (LDCT) images, reducing radiation exposure and enhancing image quality.	Significant improvements in PSNR, SSIM, and LPIPS metrics, but challenges remain in interpretability, synthetic artifacts, and clinical evaluation. Impact on precision medicine in radiology.
AI in Personalized Medicine for Disease Diagnosis and Treatment	Exploring AI's role in personalized medicine by forecasting diagnosis, disease risk, and individualized treatment plans using big data, machine learning, and deep neural networks.	AI enhances skin cancer classification, cardiology risk prediction, and chronic disease management (e.g., diabetes), improving patient outcomes, resource efficiency, and healthcare equity.

4. Future Directions and Challenges

Although the significant improvement demonstrated by GANs in low-dose computed tomography (LDCT) denoising is phenomenal, several bottlenecks stop these techniques from being fully adopted in clinical practice. Major hurdles blocking the path to implantation in actual medical settings include the absence of clinical validation, synthetic data bias, and regulatory and ethical concerns. To combat these challenges, several solutions can be considered: improving evaluation metrics, integrating hybrid models such as hybrids of GAN and other architectures, such as Diffusion Models, and increasing interpretability and explainability of GANs in the medical domain. For two interesting directions of research, one could think of combining these models with federated learning, which offers a mechanism for decentralized data processing, and the development of real-time direct application in the hospital environment, aimed at furthering the clinical applicability of GAN-based imaging solutions in precision medicine.

5. Conclusion

In conclusion, the incorporation of Artificial Intelligence (AI), specifically of Wolf Generative Adversarial Networks, in medical imaging has proven to be a potential game-changing technology in healthcare concerning diagnosis accuracy and planning of treatment, as well as the enhancement of image quality. Showing a kind of feat, GANs improve LDCT image quality, favorably reducing radiation exposure and rendering high-resolution images critical in diagnosis and treatment planning. Consequently, these rapid advances with GANs, conditional GANs (cGANs), CycleGANs, SRGANs, and some others seem to go on Ganning along the spectra of applications from medical image synthesis and domain adaptation to data augmentation and segmentation tasks. Realistic medical images very similar to clinical data have contributed to considerable progress in diagnostic imaging, and all this advancement in medical imaging is made possible through these models.

Despite these achievements, GANs still have to grapple with the multifaceted, complex issues of clinical applications. The main hurdle concerns the absence of extensive clinical validation. While it has been repeatedly shown that GANs work exceedingly well in environments controlled by the experimenter and in benchmark datasets, the question remains whether: In real clinical settings GANs will work or not. The small sizes of medical datasets mean that the clinical usability of synthetic data generated by GANs is often compromised. This adds to the

already compelling challenges of gaining acceptance for GANs to use in the clinical domain: These artefacts are at times very subtle indeed, and sometimes GAN-generated images exhibit artifactual signs of synthesis that may dissuade acceptance and use in medical diagnostics. Such challenges imply a compelling need for a more focused research effort and development with a view toward a real clinical application.

Another challenge is GAN explainability. Clinicians will demand to be assured that the results produced by AI technology can be trusted, especially in medical applications such as diagnostics. Therefore, understanding the rationale behind how GANs make predictions and ensuring that their predictions align with expectations based on valid medical standards is crucial. If interpretable and explainable, GANs will ultimately help remove the distrustful connotation surrounding AI models, allowing clinicians to have confidence in making decisions based on AI-generated data. Therefore, efforts to improve transparency concerning GAN-based systems shall be crucial to addressing these hindrances.

Ethical and regulatory issues remain prime inhibitions to AI acceptance in the medical domain. Use of simulated data and bias in generated images necessarily imply ethical consequences of trusting AI models for medical judgment. Specific guidelines regarding the permissible types of images produced by GANs should be established by regulatory agencies.

6. References

- [1] R. R. Kothinti, "Deep learning in healthcare: Transforming disease diagnosis, personalized treatment, and clinical decision-making through AI-driven innovations," 2024.
- [2] I. J. Goodfellow *et al.*, "Generative adversarial nets," *Adv Neural Inf Process Syst*, vol. 27, 2014.
- [3] A. Brock, J. Donahue, and K. Simonyan, "Large scale GAN training for high fidelity natural image synthesis," *arXiv preprint arXiv:1809.11096*, 2018.
- [4] T. Karras, S. Laine, and T. Aila, "A style-based generator architecture for generative adversarial networks," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2019, pp. 4401–4410.
- [5] J. J. Jeong, A. Tariq, T. Adejumo, H. Trivedi, J. W. Gichoya, and I. Banerjee, "Systematic review of generative adversarial networks (GANs) for medical image classification and segmentation," *J Digit Imaging*, vol. 35, no. 2, pp. 137–152, 2022.
- [6] M. Gong, S. Chen, Q. Chen, Y. Zeng, and Y. Zhang, "Generative adversarial networks in medical image processing," *Curr Pharm Des*, vol. 27, no. 15, pp. 1856–1868, 2021.
- [7] A. J. Plassard, L. T. Davis, A. T. Newton, S. M. Resnick, B. A. Landman, and C. Bermudez, "Learning implicit brain MRI manifolds with deep learning (p. 56)," 2018. [8] A. J. Plassard, L. T. Davis, A. T. Newton, S. M. Resnick, B. A. Landman, and C. Bermudez, "Learning implicit brain MRI manifolds with deep learning (p. 56)," 2018.
- [9] M. J. M. Chuquicuma, S. Hussein, J. Burt, and U. Bagci, "How to fool radiologists with generative adversarial networks? A visual turing test for lung cancer diagnosis," in *2018 IEEE 15th international symposium on biomedical imaging (ISBI 2018)*, IEEE, 2018, pp. 240–244.
- [10] H.-C. Shin *et al.*, "Medical image synthesis for data augmentation and anonymization using generative adversarial networks," in *Simulation and Synthesis in Medical Imaging: Third International Workshop, SASHIMI 2018, Held in Conjunction with MICCAI 2018, Granada, Spain, September 16, 2018, Proceedings 3*, Springer, 2018, pp. 1–11.
- [11] Y. Skandarani, N. Painchaud, P.-M. Jodoin, and A. Lalande, "On the effectiveness of GAN generated cardiac MRIs for segmentation," *arXiv preprint arXiv:2005.09026*, 2020.
- [12] S. Kazeminia *et al.*, "GANs for medical image analysis," *Artif Intell Med*, vol. 109, p. 101938, 2020.
- [13] A. Thakur and G. K. Thakur, "Developing gans for synthetic medical imaging data: Enhancing training and research," *Int. J. Adv. Multidiscip. Res*, vol. 11, no. 1, pp. 70– 82, 2024.
- [14] Z. Ju *et al.*, "HAGAN: Hybrid Augmented Generative Adversarial Network for Medical Image Synthesis," *arXiv preprint arXiv:2405.04902*, 2024.
- [15] Y. Skandarani, P.-M. Jodoin, and A. Lalande, "Gans for medical image synthesis: An empirical study," *J Imaging*, vol. 9, no. 3, p. 69, 2023.
- [16] C. Zelin, W. K. Chung, M. Jeanne, G. Zhang, and C. Weng, "Rare disease diagnosis using knowledge guided retrieval augmentation for ChatGPT," *J Biomed Inform*, vol. 157, p. 104702, 2024.
- [17] S. Pan *et al.*, "Synthetic CT generation from MRI using 3D transformer-based denoising diffusion model," *Med Phys*, vol. 51, no. 4, pp. 2538–2548, 2024.
- [18] Y. Wang, N. Yang, and J. Li, "GAN-Based Architecture for Low-dose Computed Tomography Imaging Denoising," *arXiv preprint arXiv:2411.09512*, 2024.
- [19] T. R. Gatla, "An innovative study exploring revolutionizing healthcare with ai: personalized medicine: predictive diagnostic techniques and individualized treatment," *International Journal of Advanced Research and Interdisciplinary Scientific Endeavours*, vol. 1, no. 2, pp. 61–70, 2024.